

Intro to R

Week 1

Monday

On Monday we did a brief introduction to R. We looked at making functions, variables, and played around with string substitution and printing out a little message.

```
function1 <- function(name, years){
  age <- 365*years
  final <- paste("Hello", name, "you are", age, "days old!")
  print(final)
}

function2 <- function(name, years){
  age <- 365*years
  sprintf("Hello %s, you are %d days old!", name, age)
}

function1("Ephraim", 24)
```

```
## [1] "Hello Ephraim you are 8760 days old!"
```

```
function2("Ephraim", 24)
```

```
## [1] "Hello Ephraim, you are 8760 days old!"
```

```
players <- read.csv("../data/players.csv")
games_details <- read.csv("../data/games_details.csv")
games <- read.csv("../data/games.csv")
teams <- read.csv("../data/teams.csv")
ranking <- read.csv("../data/ranking.csv")
```

Tuesday

Today we'll continue learning R and playing around with the data a little bit.

Let's start by making a function that says something about you depending on which name you put in.

```
teams$mean_arena <- mean(teams$ARENACAPACITY, na.rm = TRUE)
# hist(teams$ARENACAPACITY, breaks=12, col="red")
print(teams[teams$ARENACAPACITY == 0 , "ABBREVIATION"])
```

```
## [1] NA      NA      "ORL" NA      NA
```

We learned about tidyverse and ggplot and the powerful syntax they provide.

```
library("ggplot2")
library("tidyverse")
```

```
## -- Attaching packages ----- tidyverse 1.3.1 --
```

```
## v tibble 3.1.7      v dplyr 1.0.9
## v tidyr  1.2.0      v stringr 1.4.0
## v readr  2.1.2      v forcats 0.5.1
## v purrr  0.3.4
```

```
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()
```

We showed how things could be accomplished in both base R and in tidyverse.

```
teams$mean_size <- mean(teams$ARENACAPACITY, na.rm = TRUE)
```

```
ranking <- ranking %>% separate(STANDINGSDATE, c("year", "month", "day"))
```

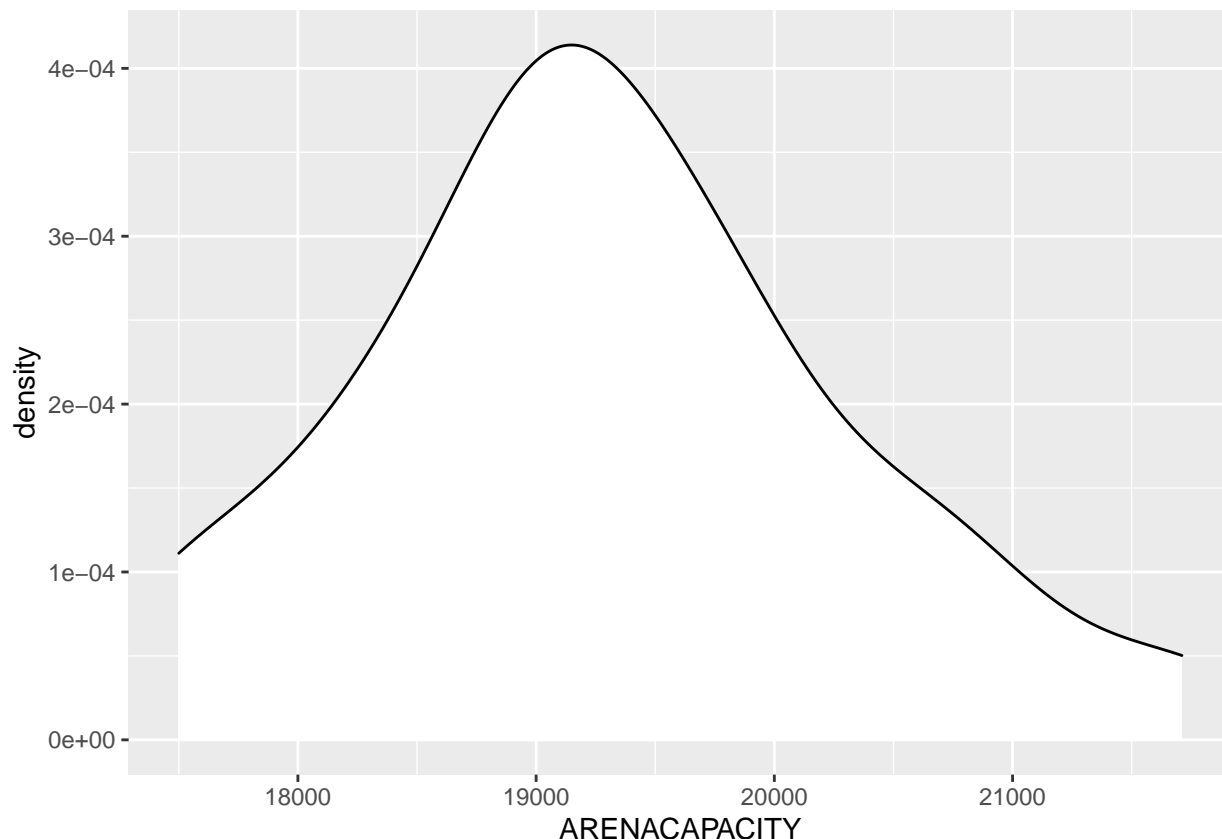
We played around with arena information and learned about missing values and adding new variables.

```
teams %>% select(NICKNAME, ARENA, ARENACAPACITY) %>% mutate(mean = mean(ARENACAPACITY, na.rm = TRUE))
```

	NICKNAME	ARENA	ARENACAPACITY	mean
## 1	Hawks	State Farm Arena	18729	18553.31
## 2	Celtics	TD Garden	18624	18553.31
## 3	Pelicans	Smoothie King Center	NA	18553.31
## 4	Bulls	United Center	21711	18553.31
## 5	Mavericks	American Airlines Center	19200	18553.31
## 6	Nuggets	Pepsi Center	19099	18553.31
## 7	Rockets	Toyota Center	18104	18553.31
## 8	Clippers	Staples Center	19060	18553.31
## 9	Lakers	Staples Center	19060	18553.31
## 10	Heat	AmericanAirlines Arena	19600	18553.31
## 11	Bucks	Fiserv Forum	17500	18553.31
## 12	Timberwolves	Target Center	19356	18553.31
## 13	Nets	Barclays Center	NA	18553.31
## 14	Knicks	Madison Square Garden	19763	18553.31
## 15	Magic	Amway Center	0	18553.31
## 16	Pacers	Bankers Life Fieldhouse	18345	18553.31
## 17	76ers	Wells Fargo Center	NA	18553.31
## 18	Suns	Talking Stick Resort Arena	NA	18553.31
## 19	Trail Blazers	Moda Center	19980	18553.31
## 20	Kings	Golden 1 Center	17500	18553.31
## 21	Spurs	AT&T Center	18694	18553.31
## 22	Thunder	Chesapeake Energy Arena	19163	18553.31
## 23	Raptors	Scotiabank Arena	19800	18553.31
## 24	Jazz	Vivint Smart Home Arena	20148	18553.31
## 25	Grizzlies	FedExForum	18119	18553.31
## 26	Wizards	Capital One Arena	20647	18553.31
## 27	Pistons	Little Caesars Arena	21000	18553.31
## 28	Hornets	Spectrum Center	19026	18553.31
## 29	Cavaliers	Quicken Loans Arena	20562	18553.31
## 30	Warriors	Chase Center	19596	18553.31

```
teams %>% filter(ARENACAPACITY != 0) %>%
  ggplot(aes(x=ARENACAPACITY)) +
  geom_density(color="black", fill="white", na.rm = TRUE, binwidth=1000)
```

```
## Warning: Ignoring unknown parameters: binwidth
```



Plot arena size

```
ranking <- ranking %>% filter(month == "03") %>% group_by(Team_ID) %>% mutate(mean_w = mean(W, na.rm = TRUE))
# ranking %>% sort(mean_w, decreasing = TRUE)

combined_games_details <- merge(games_details, games, by="GAME_ID")
```

Now we can look at the mean difference in home vs away points.

```
mean_home <- mean(games$PTS_home, na.rm = TRUE)
mean_away <- mean(games$PTS_away, na.rm = TRUE)
mean_home - mean_away
```

```
## [1] 2.811924
```

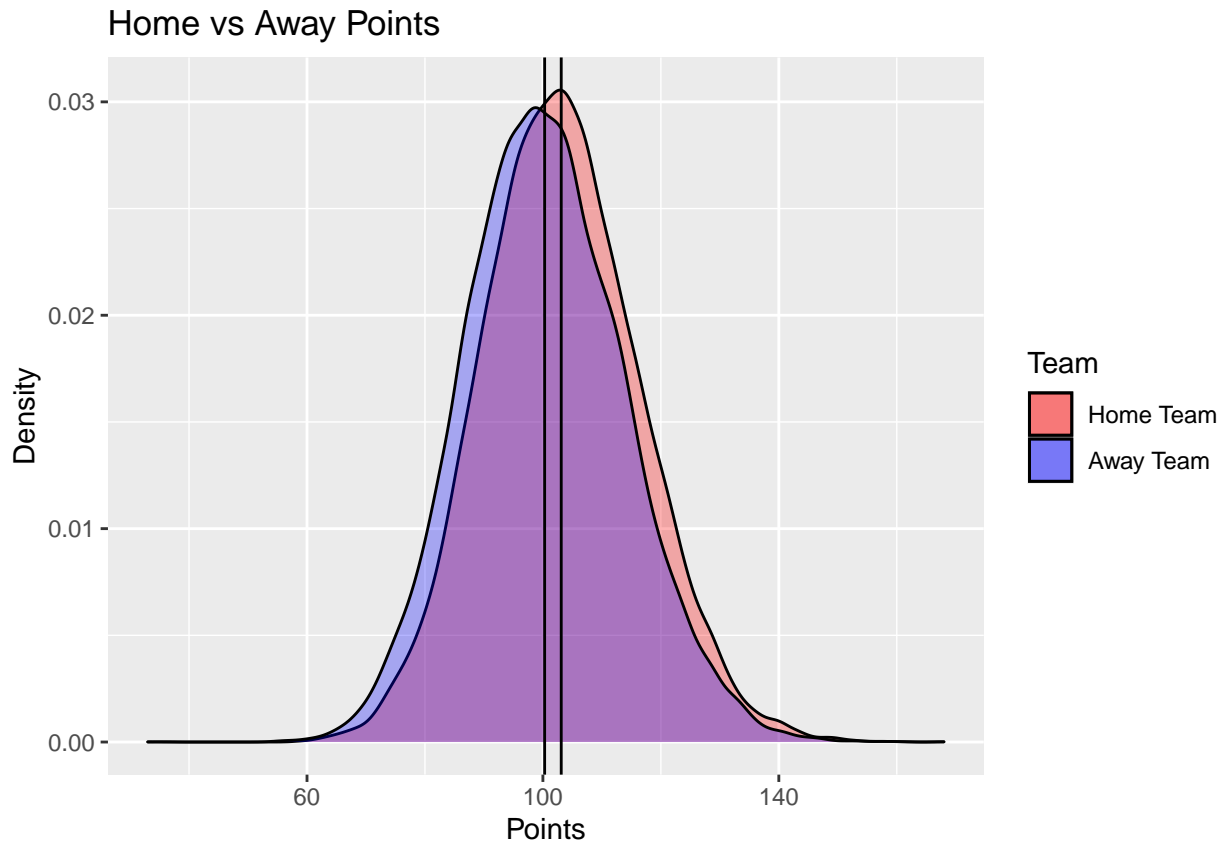
We clearly see above that the home team does marginally better. This is what many of us might expect. But the result is not that large. Do we know if this could be just by chance? In other words, what is the probability that the mean difference is instead 0?

In Statistics we have a way of testing this called a t-test. We'll come back to this.

```
ggplot() +
  geom_density(data = games, aes(x = PTS_home, fill = "home"), alpha = 0.3) +
  geom_density(data = games, aes(x = PTS_away, fill = "away"), alpha = 0.3) +
  scale_colour_manual(name = "Team", values = c("home" = "red", "away" = "blue"), labels=c("home" = "Home", "away" = "Away")) +
  scale_fill_manual(name = "Team", values = c("home" = "red", "away" = "blue"), labels=c("home" = "Home", "away" = "Away")) +
  geom_vline(xintercept=mean_home) + geom_vline(xintercept=mean_away) + labs(title="Home vs Away Points",
    x = "Points", y = "Density")
```

```
## Warning: Removed 99 rows containing non-finite values (stat_density).
```

```
## Removed 99 rows containing non-finite values (stat_density).
```



We started looking at season stats for players.

```
# games_details %>% group_by(PPLAYER_ID) %>% summarise(mean_plus_minus = mean(PLUS_MINUS, na.rm = TRUE))
```

```
games_details %>% group_by(PPLAYER_ID, PPLAYER_NAME) %>% summarise(mean_plus_minus = mean(PLUS_MINUS, na.rm = TRUE))
```

```
## `summarise()` has grouped output by 'PLAYER_ID'. You can override using the
## `.groups` argument.
```

```
## # A tibble: 2,580 x 4
## # Groups:   PLAYER_ID [2,575]
##   PLAYER_ID PLAYER_NAME    mean_plus_minus games
##   <int> <chr>          <dbl> <int>
## 1  200970 Renaldo Major      15      1
## 2  1628996 Sagaba Konate     12      1
## 3  1628431 V.J. Beachem        9      1
## 4  1629236 Jonathan Stark        9      1
## 5  1627776 Isaiah Miles          8      1
## 6  1629746 Christ Koumadje         8      1
## 7  1629166 Jeff Roberson         7.5      2
## 8    1721 Keon Clark          7      2
## 9   201939 Stephen Curry      6.56    980
## 10  203560 E.J. Singler        6.33      3
## # ... with 2,570 more rows
```

```
games_details %>% group_by(PPLAYER_ID, PPLAYER_NAME) %>% select(PPLAYER_ID, PPLAYER_NAME, PLUS_MINUS) %>% summarise(mean_plus_minus = mean(PLUS_MINUS, na.rm = TRUE))
```

```
## # A tibble: 719 x 4
## # Groups:   PLAYER_ID, PLAYER_NAME [719]
##   PLAYER_ID PLAYER_NAME      mean_plus_minus num_games
##   <int> <chr>          <dbl>      <int>
## 1    2544 LeBron James      5.16      1651
## 2    2738 Andre Iguodala     2.30      1423
## 3    2594 Kyle Korver        2.25      1409
## 4    2730 Dwight Howard      2.11      1382
## 5    2546 Carmelo Anthony     0.848     1368
## 6    1713 Vince Carter        0.889     1329
## 7   101108 Chris Paul         4.57      1309
## 8    2225 Tony Parker         4.27      1303
## 9    1717 Dirk Nowitzki       3.78      1280
## 10   2548 Dwyane Wade        2.53      1260
## # ... with 709 more rows

games_details %>% group_by(PLAYER_ID, PLAYER_NAME) %>% summarise(mean_plus_minus = mean(PLUS_MINUS, na.rm=T))

## `summarise()` has grouped output by 'PLAYER_ID'. You can override using the
## `.groups` argument.

## # A tibble: 576 x 4
## # Groups:   PLAYER_ID [576]
##   PLAYER_ID PLAYER_NAME      mean_plus_minus num_games
##   <int> <chr>          <dbl>      <int>
## 1   201939 Stephen Curry      6.56      980
## 2   202695 Kawhi Leonard       6.03      733
## 3   203110 Draymond Green      5.86      832
## 4   202691 Klay Thompson       5.80      788
## 5    1495 Tim Duncan          5.77     1128
## 6   203954 Joel Embiid        5.52      366
## 7    2544 LeBron James      5.16     1651
## 8    1938 Manu Ginobili       5.09     1199
## 9     959 Steve Nash          4.64      822
## 10  1628369 Jayson Tatum       4.62      426
## # ... with 566 more rows
```

Chef Curry

Here we played around with merging data from different dataframes.

```
players = games_details %>% select(-c("TEAM_ID", "TEAM_CITY", "PLAYER_ID", "COMMENT"))

games_date = games[,c("GAME_DATE_EST", "GAME_ID", "SEASON")]

# stats = steph.merge(games_date, on="GAME_ID", how="left")

# stats <- steph %>% left_join(games_date, by = c("GAME_ID"))

stats <- left_join(players, games_date, by = c("GAME_ID"))

head(stats)
```

```
##   GAME_ID TEAM_ABBREVIATION      PLAYER_NAME      NICKNAME START_POSITION
```

```
## 1 22101005 MIN Anthony Edwards Anthony F
## 2 22101005 MIN Jaden McDaniels Jaden F
## 3 22101005 MIN Karl-Anthony Towns Karl-Anthony C
## 4 22101005 MIN Malik Beasley Malik G
## 5 22101005 MIN D'Angelo Russell D'Angelo G
## 6 22101005 MIN Naz Reid Naz
## MIN FGM FGA FG_PCT FG3M FG3A FG3_PCT FTM FTA FT_PCT OREB DREB REB AST STL
## 1 36:22 4 10 0.400 3 8 0.375 4 4 1.00 0 8 8 5 3
## 2 23:54 6 8 0.750 1 3 0.333 1 1 1.00 2 4 6 0 0
## 3 25:17 4 9 0.444 1 3 0.333 6 8 0.75 1 9 10 0 0
## 4 30:52 4 9 0.444 4 9 0.444 0 0 0.00 0 3 3 1 1
## 5 33:46 3 13 0.231 1 6 0.167 7 7 1.00 0 6 6 9 1
## 6 23:56 3 8 0.375 1 2 0.500 4 4 1.00 3 7 10 1 3
## BLK TO PF PTS PLUS_MINUS GAME_DATE_EST SEASON
## 1 1 1 1 15 5 2022-03-12 2021
## 2 2 2 6 14 10 2022-03-12 2021
## 3 0 3 4 15 14 2022-03-12 2021
## 4 0 1 4 12 20 2022-03-12 2021
## 5 0 5 0 14 17 2022-03-12 2021
## 6 2 1 1 11 -7 2022-03-12 2021
```

We then learned about how to chain commands together in the tidyverse (or dplyr) syntax using pipes. We learned how to group datasets by variables of interest and perform computations within these groups.

```
seasonal_stats <- stats %>% group_by(SEASON, PLAYER_NAME) %>% summarise(PTS = mean(PTS, na.rm=TRUE)) %>%
```

```
## `summarise()` has grouped output by 'SEASON'. You can override using the
## `.groups` argument.
```

```
seasonal_stats
```


```
## # A tibble: 10,707 x 3
## # Groups:   SEASON [19]
## SEASON PLAYER_NAME PTS
## <int> <chr> <dbl>
## 1 2003 A.J. Guyton 4
## 2 2003 Aaron McKie 9.16
## 3 2003 Aaron Williams 5.89
## 4 2003 Ademola Okulaja 2
## 5 2003 Adonal Foyle 3.11
## 6 2003 Adrian Griffin 0.579
## 7 2003 Al Harrington 12.6
## 8 2003 Alan Henderson 4
## 9 2003 Alex Garcia 1.5
## 10 2003 Alex Scales 13.5
## # ... with 10,697 more rows
```

Last weekends homework.

```
select_players <- function (players) {
  seasonal_stats %>% filter(PLAYER_NAME %in% players) %>%
  ggplot(aes(x=factor(SEASON), y=PTS, group=PLAYER_NAME, fill=PLAYER_NAME, colour=PLAYER_NAME)) + geom_line()
  theme( axis.text.x=element_text(angle=90) )
}
```

```
select_players(players=c("Stephen Curry", "LeBron James", "Jayson Tatum", "Kawhi Leonard"))
```

```
## Warning: Removed 1 row(s) containing missing values (geom_path).
```



Curry vs the rest of the league

Here we learned how to summarise (collapse the data) across multiple columns or variables at once. We then plotted performance for the rest of the team against a player of our choice.

```
all_players <- games_details %>% select(PLAYER_NAME, FGM, FG_PCT, FG3_PCT, PTS, FG3M, FG3A, FTM, FT_PCT)
ggplot() + geom_point(data=all_players, aes(x=FGM, y=PTS)) + geom_point(data=all_players[all_players$PL
```

Steph Curry vs the league
Points vs Number of Fieldgoals Made

